

Distributed snapshots

- Global states of the system that are consistent with causality
- Such global states can be used for detecting **stable properties**, e.g.
 - Distributed deadlock detection
 - Distributed garbage collection
 - Distributed termination detection
- Main difficulty
 - System state changes during observation

Distributed snapshots: assumptions

- The system is connected, that is there is a path between every pair of processors
- $C_{i,j}$ channel from p_i to p_j
- Channels are reliable and FIFO
 - Messages sent are eventually received in order

Distributed snapshots: definitions

- Processor state
 - State of a processor (at an instant) is the assignment of a value to each variable of that processor
- Channel state
 - State of $C_{i,j}$ is the ordered list of messages sent by p_i but not yet received at p_j


Distributed snapshots: definitions

- Global state of the system
 - A pair (S, L) where
 - $S = (s_1, \dots, s_M)$ denotes the processor states and
 - L = channel states
- Note
 - A global state cannot be taken instantaneously
 - It must be computed in a distributed manner

Distributed snapshots

- The problem
 - Devise a distributed algorithm that computes a *consistent global state*.
- What do we mean by consistent global state?

Distributed snapshot: meaning of consistent global state



Two possible states for each processor: s_0, s_1


In s_0 : the processor hasn't the token

In s_1 : the processor has the token

The system contains exactly one token which moves back and forth between p and q . Initially, p has the token.

Events: sending/receiving the token


Distributed snapshots: meaning of consistent global state



Distributed snapshots: meaning of consistent global state

A global state G is consistent if it is one that **could have occurred**

Consider a system with two possible runs (non-determinism!)



—→ Actual transitions

The output of the snapshot algorithm can be G !

Distributed snapshots

- $S = \{s_1, \dots, s_M\}$
- o_i : event of observing s_i at p_i
- $O(S) = \{o_1, \dots, o_M\}$
- Definition
 - S is a **consistent cut** iff $O(S) = \{o_1, \dots, o_M\}$ is consistent with causality


Distributed snapshots

- Definition:


- $\{o_1, \dots, o_M\}$ is consistent with causality iff
$$(\forall e : e \text{ in } E_i \wedge e <_H o_i : \\ (\forall e' : e' \text{ in } E_j \wedge e' <_H e : e' <_H o_j))$$

- Such an observation cannot indicate the receipt of a message without the sending of that message

Distributed snapshots: intuition




OK




KO

Distributed snapshots

Exercise



Is $O=\{o_1, o_2, o_3\}$ consistent with causality?




Is $O'=\{o'_1, o'_2, o'_3\}$ consistent with causality?

Distributed snapshots: messages sent but not yet received

- Let m be a message, we write
 - $s(m)$: sending event of message m
 - $r(m)$: receipt event of message m
- Let S be an observed global state
 - $O(S) = \{o_1, o_2, \dots, o_M\}$
 - If $s(m) <_{p_i} o_i$ and $o_j <_{p_j} r(m)$ then message m is sent but not received relatively to observation $O(S)$

Distributed snapshots: messages sent but not yet received



Processor q observes its state at **1**
Then, q sends **special message** to p
asking p to record its state

Processor p records its state at **2**
then p sends a **special message** to q
saying that it recorded its state

When q receives the **special message**
from p , processor q determines that
messages m_1 , m_2 , m_3 were sent but
were not yet received

In the resulting global state, the state
of C_{pq} must be **$[m_1, m_2, m_3]$**

Distributed snapshots

- Use of non-consistent global state can lead to wrong conclusions
- E.g. a false deadlock can be claimed!


False deadlock illustrated

- Assume
 - r_1 and r_2 two resources
 - Each resource is either *available* or *un-available*
 - p_1 and p_2 two processors that access resources
- To access a resource r , a processor p sends Request (Req) to r and p starts to wait until it receives Ok from r

False deadlock illustrated

- When a resource **r** in status *available* receives Req from **p**,
 - **r** sends Ok to **p**
 - **r** becomes *un-available*
 - **r** starts to wait until it receives Release (Rel) from **p**
- A deadlock situation occurs when there is a cycle in the wait-for-graph amongst processors and resources

False deadlock illustrated




Inference from

o_1 : p_1 waits for r_2

o_2 : r_1 waits for p_1

o_3 : r_2 waits for p_2

o_4 : p_2 waits for r_1



From $O = \{o_1, o_2, o_3, o_4\}$
the deadlock detector
concludes there is a
deadlock!

What's wrong with O ?

Distributed snapshots: the algorithm

- For simplicity
 - We assume only one snapshot computation
- Principle
 - Uses two types of messages
 - **SnapshotRequest**
 - Serve to start a snapshot computation
 - **SnapshotToken**
 - Serve to “propagate” a snapshot computation
 - Each processor records its state and the state of its incoming channels

Distributed snapshots: the algorithm

- A processor that receives **SnapshotRequest**
 - Records its state
 - Sends a **SnapshotToken** to each of its outgoing channels
 - Starts to record state of incoming channels
- Recording of the state of an incoming channel is finished when a **SnapshotToken** is received along it

Distributed snapshots: the algorithm

- A processor p that receives **SnapshotToken** on channel C_{qp}
 - If this is the first time that **SnapshotToken** is received then
 - p records its state
 - p records the state of C_{qp} as empty
 - p sends one **SnapshotToken** to each outgoing channel
 - p starts to record the state of all its other incoming channels
 - Otherwise, p records the state of C_{qp} as the list of messages received from q since the time p recorded its state

Distributed snapshots: the algorithm

- Local termination of a snapshot
 - A processor learns that its participation to a snapshot computation is terminated when it has recorded its state and received a **SnapshotToken** on each of its incoming channels
- When does an initiating processor learn that the global snapshot computation is terminated?
 - We'll see later

Distributed snapshots: the algorithm (pseudo-code)

- Constants

- Incoming: set of processors from which I receives
- Outgoing: set of processors that receive from me

- Variables

- MyVersion: Integer
- StateForSnapshotOf: $\{1, \dots, M\} \rightarrow \text{States}$
- CurrentSnapshotVersionOf: $\{1, \dots, M\} \rightarrow \text{Integer}$
- TokensReceivedForSnapshotOf: $\{1, \dots, M\} \rightarrow \text{Integer}$
- ChannelState: $\{1, \dots, M\} \rightarrow (\{1, \dots, M\} \rightarrow \text{MessageLists})$

Distributed snapshots: the algorithm (pseudo-code)

- Variable initialization
 - $\text{MyVersion} := 0$
 - $\text{StateForSnapshotOf}(i) := ?$, $1 \leq i \leq M$
 - $\text{CurrentSnapshotVersionOf}(i) := 0$, $1 \leq i \leq M$
 - $\text{TokensReceivedForSnapshotOf}(i) := ?$, $1 \leq i \leq M$
 - $(\text{ChannelState}(i))(j) := ?$, $1 \leq i, j \leq M$
- $?$: we don't care!

Distributed snapshots: the algorithm (pseudo-code)

```
Wait for SnapshotRequest or SnapshotToken(r,v)  
on SnapshotRequest do  
    MyVersion := MyVersion + 1  
    StateForSnapshotOf(self) := relevant local state  
    CurrentSnapshotVersionOf(self) := MyVersion  
    for every q in Outgoing do  
        send(q, SnapshotToken(self, MyVersion))  
    od  
    TokenReceivedForSnapshotOf(self) := 0  
od  
...
```

Distributed snapshots: the algorithm (pseudo-code)

Wait for **SnapshotRequest** or **SnapshotToken(r,v)**

...

on **SnapshotToken(r,v)** from q do

if CurrentSnapshotVersionOf(r) < v then

StateForSnapshotOf(r) := relevant local state

CurrentSnapshotVersionOf(r) := v

(ChannelState(r))(q) := **nil**

for every k in Outgoing do

send(k, SnapshotToken(r,v))

od

TokenReceivedForSnapshotOf(r) := 1

/* prepare for incoming channel state computation*/

else

...

od

Distributed snapshots: the algorithm (pseudo-code)

Wait for **SnapshotRequest** or **SnapshotToken(r,v)**

...

on **SnapshotToken(r,v)** from q do

if CurrentSnapshotVersionOf(r) < v then

 ...

else

if CurrentSnapshotVersionOf(r) = v then

 TokensReceivedForSnapshotOf(r) := TokensReceivedForSnapshotOf(r) + 1

 (ChannelState(r))(q) := list of msg received from q since I recorded my state

if TokensReceivedForSnapshotOf(r) = | Incoming | then

 my local participation is finished for snapshot computation (r,v)

od

Distributed snapshots

the algorithm's properties

- Shows an important technique for designing distributed algorithms: **Diffusing computation**
 - The computation starts at one processor, then the computation diffuses through the whole set of processors
- Shows how to distinguish between objects that have same name, but actually different
 - The use of **version numbers**
- Shows how to flush communication channels
 - The use of **SnapshotToken** messages
- Can record a global state that never occurred

Distributed snapshots: adding termination detection

- The algorithm for the snapshot computation can be easily adapted to construct a **spanning tree** of the system
- Spanning tree of the system
 - Each node maintains a variable MyParent [1..M]
 - A processor that starts a snapshot computation is the root of the associated spanning tree
 - The processor from which I receive **SnapshotToken(r,v)** for the first time is my parent in spanning tree associated to the snapshot computation identified by (r,v)

Distributed snapshots: adding termination detection

- Using the spanning tree associated to a snapshot computation, we can let the processor that started that snapshot computation detect when the computation is terminated

Distributed snapshots: adding termination detection

- **Idea**

- When I receive a SnapshotToken(r, v) for the first time, I send SnapshotToken(r, v) to all outgoing channels except to the one leading to MyParent(r)
- When I have received SnapshotToken(r, v) on all my incoming channels
 - If $r \neq \text{self}$ then I send SnapshotToken(r, v) to MyParent(r)
 - Otherwise, I learn that the global computation is finished

Distributed snapshots: adding termination detection

```
Wait for SnapshotRequest or SnapshotToken(r,v)  
on SnapshotRequest do  
    MyVersion := MyVersion + 1  
    StateForSnapshotOf(self) := relevant local state  
    CurrentSnapshotVersionOf(self) := MyVersion  
    for every q in Outgoing do  
        send(q, SnapshotToken(self, MyVersion))  
    od  
    TokenReceivedForSnapshotOf(self) := 0  
    MyParent(self) := self  
od  
...
```


Distributed snapshots: adding termination detection

Wait for **SnapshotRequest** or **SnapshotToken(r,v)**

...

on **SnapshotToken(r,v)** from q do

if **CurrentSnapshotVersionOf(r)** < v then

MyParent(r) := q

StateForSnapshotOf(r) := relevant local state

CurrentSnapshotVersionOf(r) := v

 (**ChannelState(r)**)(q) := **nil**

for every k in **Outgoing** and $k \neq C_{\text{self},q}$ do

send(k, **SnapshotToken(r,v)**)

od

TokenReceivedForSnapshotOf(self) := 1

 /* prepare for incoming channel state computation*/

else

 ...

od

Distributed snapshots: adding termination detection

```
Wait for SnapshotRequest or SnapshotToken(r,v)  
  on SnapshotToken(r,v) from q do  
    if CurrentSnapshotVersionOf(r) < v then  
      ...  
    else  
      if CurrentSnapshotVersionOf(r) = v then  
        TokensReceivedForSnapshotOf(r) := TokensReceivedForSnapshotOf(r) + 1  
        (ChannelState(r))(q) := list of msg received from q since I recorded my state  
        if TokensReceivedForSnapshotOf(r) = | Incoming | then  
          if r = self then  
            global termination  
          else  
            send(MyParent(r), SnapshotToken(r,v))  
      od
```

Distributed snapshots

- Discussion
 - Concurrent observations (see book)
- Exercise
 - Explain why the proposed algorithm computes consistent global states

Modeling a distributed computation

- How to prove that a given distributed algorithm is correct?
- A formal model of a distributed computation serves this purpose
- **Idea**
 - Model a distributed computation as a sequence of global state transitions where each global state transition is caused by the execution of one event

Modeling a distributed computation

- Initial state of the system
 - $\mathbf{G}^0 = (\mathbf{S}^0, \mathbf{L}^0)$ where
 - \mathbf{S}^0 is the vector of processor initial states
 - \mathbf{L}^0 is the channel initial states, each channel is empty
- Given a global state \mathbf{G} , a processor \mathbf{p}
 - We write $\mathbf{s}_p|_{\mathbf{G}}$ for state of \mathbf{p} in global state \mathbf{G}
 - We write $\mathbf{State}(\mathbf{c})|_{\mathbf{G}}$ for state of channel \mathbf{c} in \mathbf{G}

Modeling a distributed computation

- Events

- An event \mathbf{e} is a 5-tuple $(\mathbf{p}, \mathbf{s}, \mathbf{s}', \mathbf{m}, \mathbf{c})$ where
 - \mathbf{p} is a processor
 - \mathbf{s}, \mathbf{s}' are possible states of \mathbf{p}
 - \mathbf{m} is a message or Null
 - \mathbf{c} is a channel or Null
- Interpretation of an event $\mathbf{e} = (\mathbf{p}, \mathbf{s}, \mathbf{s}', \mathbf{m}, \mathbf{c})$
 - Takes \mathbf{p} from \mathbf{s} to \mathbf{s}'
 - Possibly sends or receives \mathbf{m} on \mathbf{c}

Modeling a distributed computation

- Let
 - $\mathbf{G}=(\mathbf{S},\mathbf{L})$ be a global state
 - $\mathbf{e}=(p, \mathbf{s}, \mathbf{s}', \mathbf{m}, \mathbf{c})$ an event
- Event \mathbf{e} *can occur* in \mathbf{G} if
 - There is a global state \mathbf{G}' such that executing \mathbf{e} in \mathbf{G} changes the system state from \mathbf{G} to \mathbf{G}'
 - **Notation:** $\mathbf{G} \rightarrow_{\mathbf{e}} \mathbf{G}'$

Modeling a distributed computation

- If $G \rightarrow_e G'$ then G' is defined by
 - $s_p|_{G'} = s'$
 - For every $q \neq p$, $s_q|_{G'} = s_q$
 - If c is an outgoing channel from p , then
 - $\text{State}(c)|_G = [m_1, m_2, m_3, \dots, m_k, m]$
 - If c is an incoming channel to p and $\text{State}(c)|_G = [m, m_1, m_2, \dots, m_k]$ then
 - $\text{State}(c)|_{G'} = [m_1, m_2, m_3, \dots, m_k, m]$
 - For every $c' \neq c$, $\text{State}(c')|_{G'} = \text{State}(c')|_G$
- **Note**
 - If $c = \text{Null}$, then for every c' , $\text{State}(c')|_{G'} = \text{State}(c')|_G$

Modeling a distributed computation

- Let
 - \mathbf{G} be a global state
 - $\mathbf{seq} = \langle \mathbf{e}_0, \mathbf{e}_1, \dots, \mathbf{e}_j \rangle$ be a sequence of events
- \mathbf{seq} *can occur* in \mathbf{G} if there is a sequence of global states $\langle \mathbf{G} = \mathbf{G}_0, \mathbf{G}_1, \dots, \mathbf{G}_{j+1} = \mathbf{G}' \rangle$ such that $\mathbf{G}_i \rightarrow_{\mathbf{e}_i} \mathbf{G}_{i+1}$ for $i=0, \dots, j$
 - **Notation:** $\mathbf{G} \rightarrow_{\mathbf{seq}} \mathbf{G}'$

Modeling a distributed computation

- Let \mathbf{G} , \mathbf{G}' be two global states
- We say that \mathbf{G}' is reachable from \mathbf{G} if there is a sequence of events \mathbf{seq} such that $\mathbf{G} \rightarrow_{\mathbf{seq}} \mathbf{G}'$
- We define a computation of the system as a sequence of events \mathbf{seq} such that
 - $\mathbf{G}^0 \rightarrow_{\mathbf{seq}} \mathbf{G}'$ for some \mathbf{G}'

Modeling a distributed computation

- Characteristics of this model
 - Each global state transition is caused by the execution of one event that can occur in the current global state
- But several events can occur at a given global state
 - In this case, one of them is randomly selected
 - i.e. non-deterministic computations

Modeling a distributed computation

- How to prove that a given distributed algorithm is correct in this model
 - We must show that **every** global state reachable from the initial global state satisfies the correctness properties
- **Stable properties**
 - A global state predicate Φ such that if $\Phi(G)$ is true then $\Phi(G')$ is true for every G' reachable from G

Relationships between models

- So far, three models for a distributed computation that can be grouped into two classes
- **Observational models**
 - Model 1
 - Set of events related by Happens-before relation (H-DAG)
 - Model 2
 - Consistent cuts
 - Framework for analyzing what happened in an execution
 - Concurrency of events is explicitly captured

Relationships between models

- Predictive model

- Model 3

- Global state transitions ([interleaving](#))
 - Given an initial global state \mathbf{G}^0 and a distributed algorithm \mathbf{A} , one can say what might occur when \mathbf{A} is executed starting from \mathbf{G}^0
 - Captures consistent cuts and causality
 - Concurrency is not explicitly captured, instead non-determinism is used

Relationships between models

- Given
 - **A** a distributed algorithm
 - **H-DAG(A)**, an execution of **A**
 - **E** all events in **H-DAG(A)**
 - **Z(E)** all consistent cuts of **H-DAG(A)**
- There is an **execution DAG**, **G(E)** that corresponds to **H-DAG(A)**
 - **G(E)** shows all possible executions, in the Model 3, that can be derived from **H-DAG(A)**

Relationships between models

- Construction of $G(E)$
 - Since events in each processor are totally ordered, they can be numbered sequentially
 - Let **Number(e)** be the number assigned to event **e**


Relationships between models

- Construction of $G(E)$
 - Let $z_1 = (z_{11}, z_{12}, \dots, z_{1M})$ and $z_2 = (z_{21}, z_{22}, \dots, z_{2M})$ be two consistent cuts
 - $z_1 <_H z_2$ if
 - $\text{Number}(z_{1j}) \leq \text{Number}(z_{2j})$ for $j=1, 2, \dots, M$
 - z_2 *immediately follows* z_1 if there is a j such that
 - $\text{Number}(z_{2k}) = \text{Number}(z_{1k})$ for $k=1, 2, \dots, M$ and $k \neq j$
 - $\text{Number}(z_{2j}) = \text{Number}(z_{1j}) + 1$


Relationships between models

- Construction of $G(E)$
 - $G(E)$ is a **directed graph** where
 - The nodes are elements of $Z(E)$
 - There is an arc from z_1 to z_2 if z_2 immediately follows z_1
 - Each node $z=(z_1, \dots, z_M)$ of $G(E)$ is labeled by a vector (n_1, n_2, \dots, n_M) where n_i is the number of z_i
 - Each arc in of $G(E)$ represents the execution of one event (a global state transition)

Execution DAG illustrated



H-DAG



Execution-DAG

Execution DAG predicates

- Predicates on execution DAG
- Useful for
 - Expressing desirable properties
 - Debugging
 - Proving correctness

Execution DAG predicates

- Let
 - $G(E)$ be an execution DAG
 - Φ be a global state predicate
 - z be a node of $G(E)$
- We define **Concur(z)** as the set of all nodes z' of $G(E)$ such that there is no path from z to z' or from z' to z
 - E.g. **112** and **131** of the example execution DAG

Execution DAG predicates

- **POSSIBLY** predicate
 - $\text{POSSIBLY}(\Phi, G(E))$ is true if
 - There is a node z of $G(E)$ such that $\Phi(z)$ is true
- **ALLPATH** predicate
 - $\text{ALLPATH}(\Phi, G(E))$ is true if
 - For every path P from the source to the sink of $G(E)$, there is z in P such that $\Phi(z)$ is true

Execution DAG predicates

- **DEFINITELY** predicate
 - $\text{DEFINITELY}(\Phi, G(E))$ is true if
 - There is a node z of $G(E)$ such that $\Phi(z')$ is true for every z' in $\text{Concur}(z)$
- **ALWAYS** predicate
 - $\text{ALWAYS}(\Phi, G(E))$ is true if
 - For every node z of $G(E)$, $\Phi(z)$ is true

Summarizing future executions

- Ways to describe future states of an execution given the current state
- Definition
 - Let G_0 be a global state. A **run** r on G_0 is defined as a finite or an infinite sequence of events $\langle e_0, e_1, e_2, \dots \rangle$ that *can occur* in G_0
- Notation:
 - $r[i]$: sequence formed by the first $i+1$ events of r , $0 \leq i \leq |r|-1$
 - $R(G_0)$ denotes all possible runs on G_0

Summarizing future executions

- *Eventually* predicate
 - Let Φ be a global state predicate
 - Let Ψ be a predicate on $\mathbf{R}(G_0)$
 - $\text{Eventually}(\Phi, G_0, \Psi)$ is true iff
 - For every \mathbf{r} in $\mathbf{R}(G_0)$ such that $\Psi(\mathbf{r})$ is true,
 - There exists i in $\{0, 1, 2, \dots, |\mathbf{r}|-1\}$
 - There exists a global state G such that $G_0 \rightarrow_{\mathbf{r}[i]} G$ and $\Phi(G)$ is true

Summarizing future executions

- *Always* predicate
 - Let Φ be a global state predicate
 - Let Ψ be a predicate on $\mathbf{R}(G_0)$
 - $Always(\Phi, G_0, \Psi)$ is true iff
 - For every \mathbf{r} in $\mathbf{R}(G_0)$ such that $\Psi(\mathbf{r})$ is true,
 - For every i in $\{0, 1, 2, \dots, |\mathbf{r}|-1\}$
 - There exists a global state G such that $G_0 \rightarrow_{\mathbf{r}[i]} G$ and $\Phi(G)$ is true

Failures in a distributed system

- A distributed system like any other computing system is built on
 - Hardware and software components that can fail from time to time
- A faulty component can behave in an unpredictable manner

Failures in a distributed system

- But, in many cases we must design distributed systems in a way that enables them to achieve their well-defined goals despite the presence of failure of some of its components
 - Hospital
 - Banking
 - Transportation, ...
- The system must be fault-tolerant

Failures in a distributed system

- Fault-tolerance requires
 - Failure semantics
 - Description of the way a faulty component behaves
 - Mechanisms/techniques for handling well-defined failures
- Any interaction between two components of a distributed system can be seen as involving
 - A server: provides service to clients
 - A client: requests service from the server

Failures in a distributed system: failure semantics (F. Cristian 91)

- Omission failure
 - Occurs when a server omits to respond to a request
- Response failure
 - Occurs when a server responds incorrectly to a request
 - Wrong state transition at the server
 - Wrong value is returned to the client
- Timing failure
 - Occurs when a server responds either too earlier or too late

Failures in a distributed system: failure semantics

- Crash failure

- Occurs when a server repeatedly omits to respond to requests until it is **restarted**
- **Cannot be accurately detected**
- Depending upon the state of the server when it restarts (if any restart), several subclasses of crash failure are given

Failures in a distributed system: failure semantics

- Amnesia-crash
 - Crash failure and
 - In addition, when the server restarts, it restarts from its initial state, the state by the time of failure is lost
- Pause-crash
 - Crash failure and
 - In addition, when the server restarts, it restarts from the state before the crash

Failures in a distributed system: failure semantics

- Halting-crash
 - Crash failure and
 - The server never restart
- Fail-stop
 - A fail-stop server is one that either functions correctly or stops. And when stopped this fact is accurately detected by its clients
 - Very nice but not realistic

Failures in a distributed system: failure semantics

- Byzantine failure

- A Byzantine server can exhibit any of
 - Crash, Timing, Response and Omission failure
- A Byzantine server can send contradictory messages to different recipients
- The most general failure semantics, captures malicious behaviors
- Often assumed for ultra-reliable systems